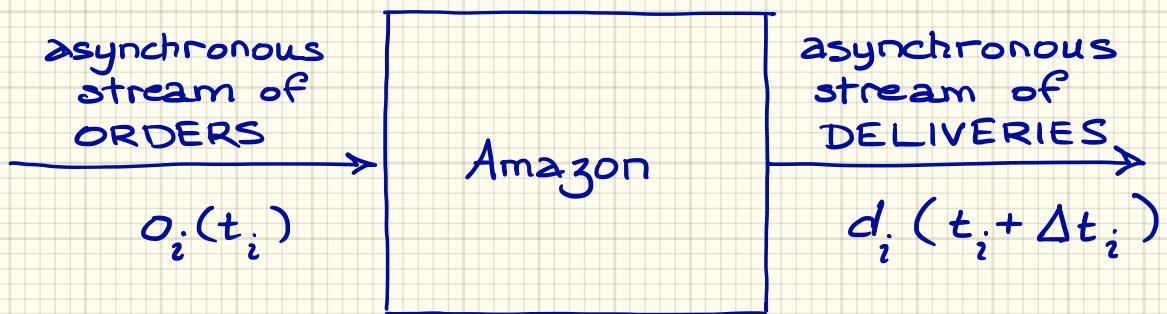


Time Warp 2018

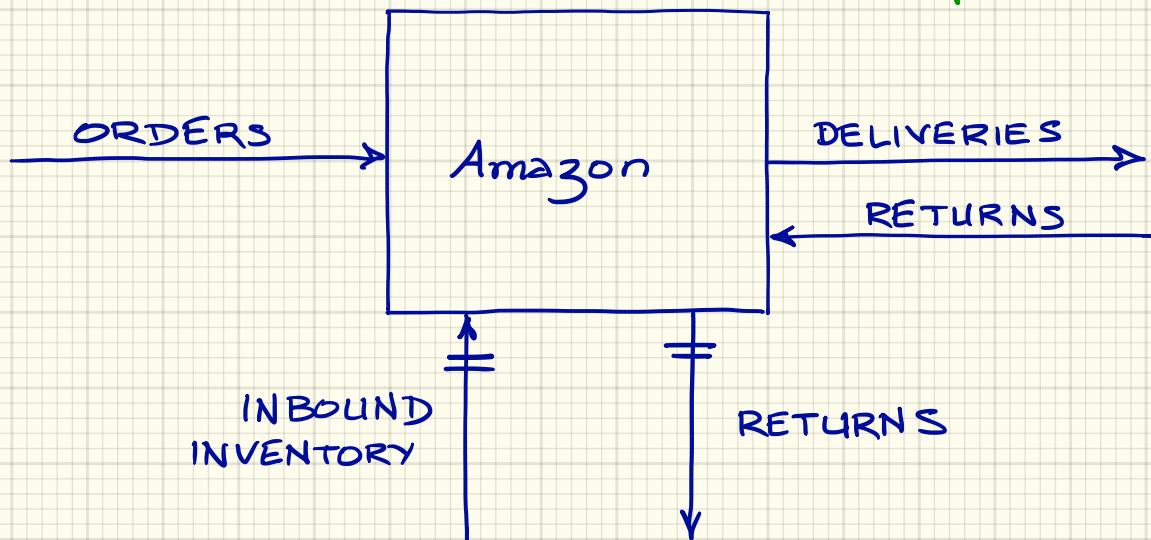
Distributed Discrete-Event  
Simulation

Amazon is a Function from  
orders to deliveries

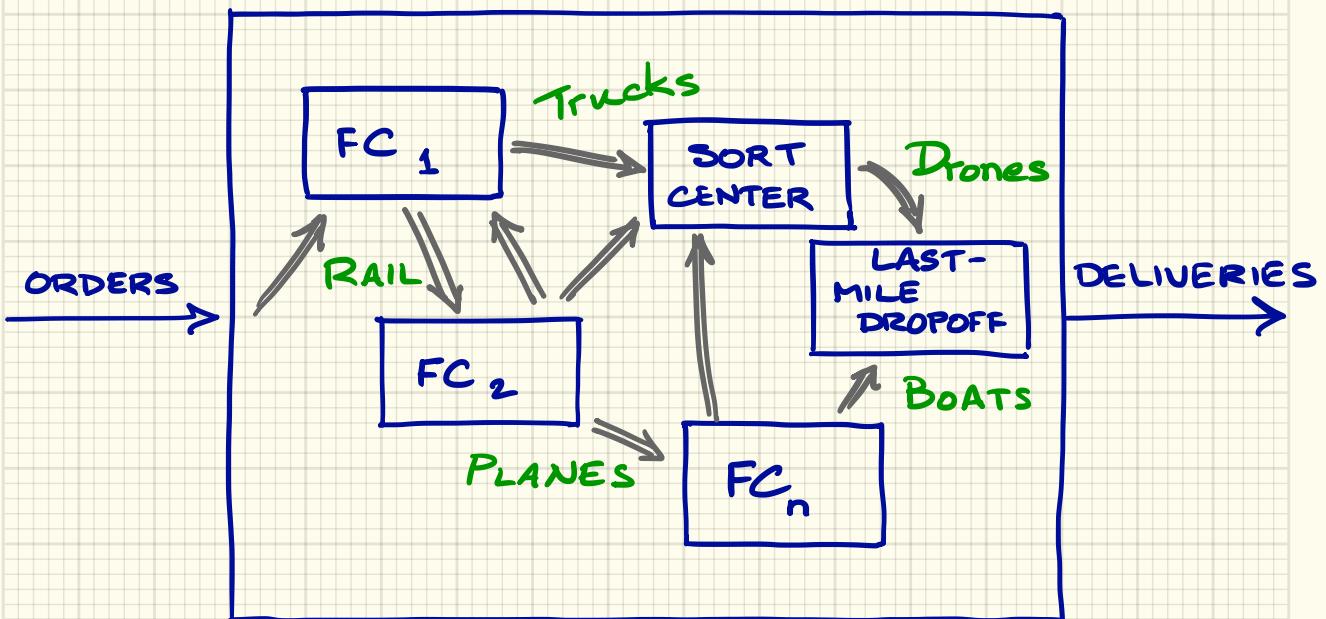


The function itself depends on time due to inventory & returns

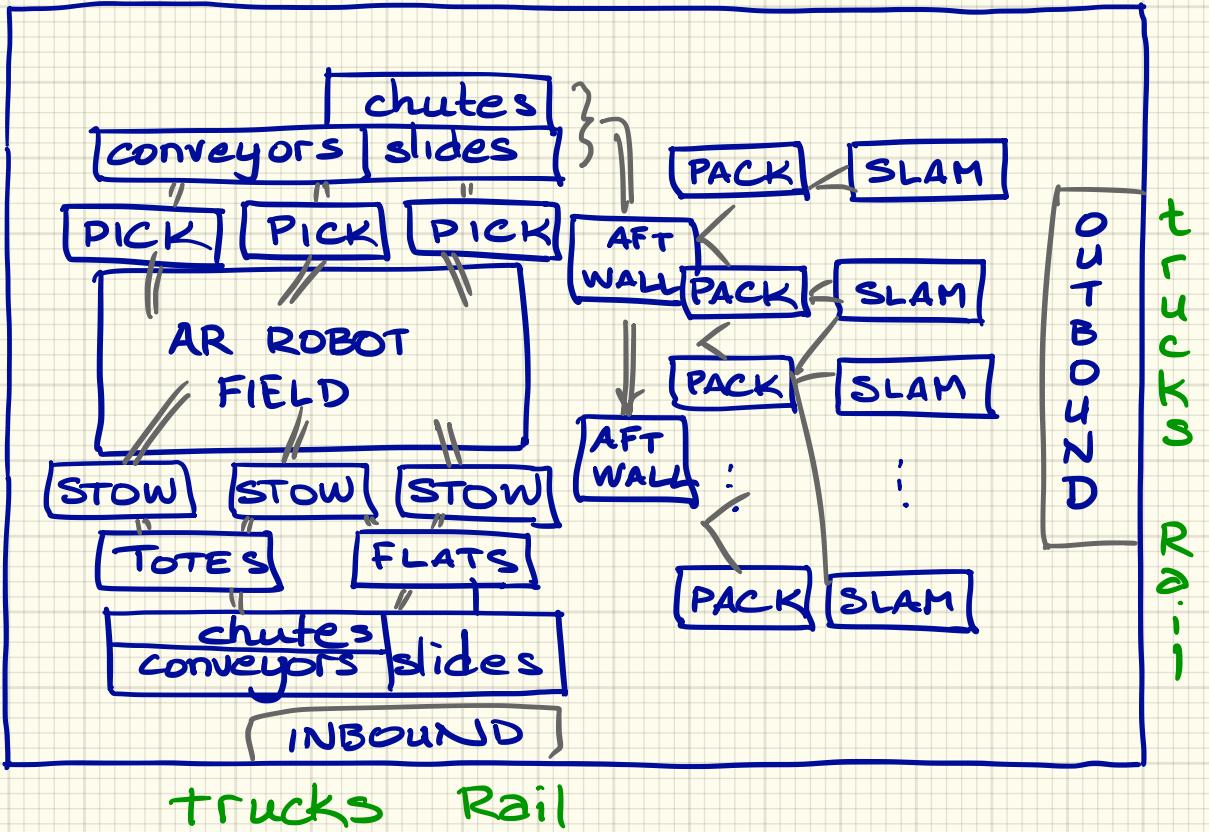
we'll ignore that for now



# ONE LEVEL Down



# ONE MORE LEVEL DOWN



## Questions:

- **geography & velocity of demand**

⇒ how many FC's? where?

⇒ how many conveyors? robots? in  
the FC, sort center, last-mile station

⇒ how many trucks, planes, containers?  
where? when?

- **FINANCE** is a whole 'nuther  
dimension ...

⇒ how much? borrow?  
sell stock? when?

Billion-  
dollar  
decisions

maybe even  
compose with  
detailed mech  
sims of indivi-  
dual robots &  
vehicles

Analytic solutions out of the question

Queuing model with fixed flows?

~~OR~~

Discrete-Event Simulation

- set up components & flows
- do "what-if's" (COUNTERFACTUALS)
  - put a big FC in { Kansas, Texas, Alaska }
  - run some sims, see what happens
  - put { more / fewer } { faster / slower }  
{ conveyors / transports }  
see what happens

# Ideal for Time Warp

Millions of logical processes  
irregular space & time structure  
runs much too slowly NEED SPEEDUP

fast things

planes  
conveyors  
orders

big things

fc's  
trucks  
rail cars

slow things

trucks  
trains  
AR fields

small things

Asins  
boxes  
totes

Time - stepped sim : step every process at  
the rate of the slowest one  
⇒ most processes idle most of the time

Optimistic sim (aka speculative) : processes  
run with best-guess inputs, rollback  
& cancel incorrect outputs when  
necessary

Time - stepped sim : step every process at  
the rate of the slowest one  
⇒ most processes idle most of the time

Optimistic sim (aka speculative) : processes  
run with best-guess inputs, rollback  
& cancel incorrect outputs when  
necessary CAN WE DO THIS DISTRIBUTED,  
IN PARALLEL, GET SPEEDUP  
FROM PARALLEL EXECUTION ?



Synchronous, Conservative, parallel simulation has been studied extensively:

Never do incorrect computation

Block & wait or deadlock - and break

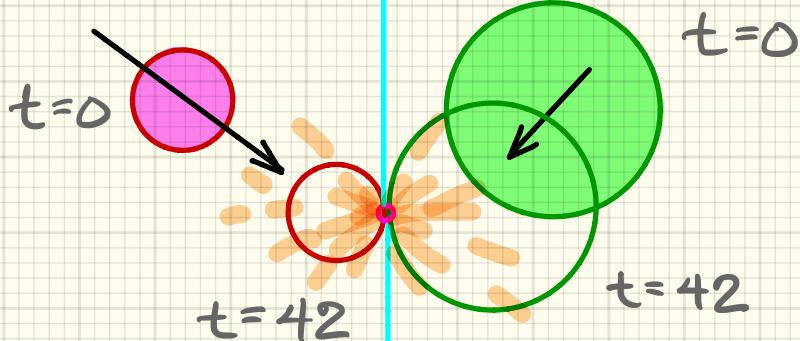
Requires order-preserving queues

Lacks essential symmetries of TW

Requires static graph of comm. channels

Slower than TW for irregular problems

(faster than TW for some network sims)



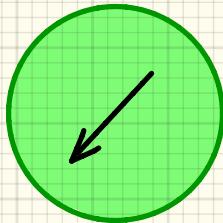
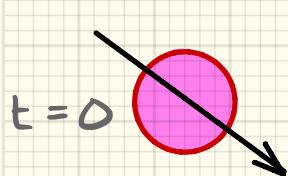
**COLLISION  
Predicted**

- FOR time 42
- AT time 0

**DETAILS OF  
THE DEMO**

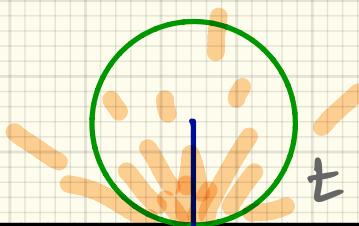
**PROGRAMMING  
MODEL**

sender : small puck  
 receiver : big puck  
 send time : 0  
 receive time : 42  
 body : your new  
 center & velocity  
 will be ...



$t = 0$

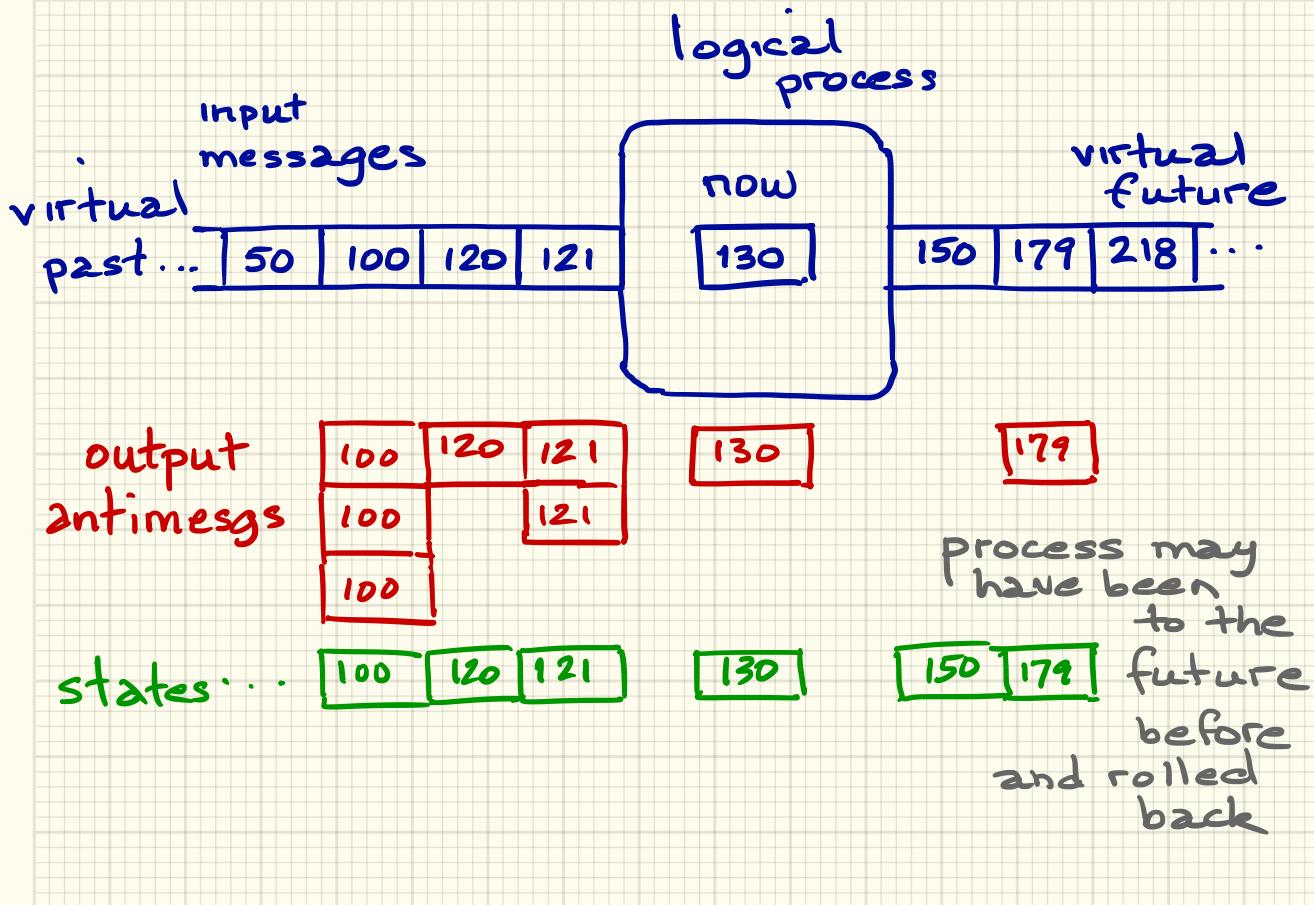
But, but . . .



$t = 79$

but big puck didn't know  
about small & already  
scheduled a later collision  
with the wall !

what to do? : small causes big to roll back &  
cancel the incorrect wall collision by sending  
antimessage to self. SYSTEM DOES THIS  
AUTOMATICALLY — NO USER CODE NEEDED



## Side Note:

All forms of speculative computation at all levels can benefit from Time Warp

{Synchronization by rollback alone}  
{Cancellation by antimesage alone}

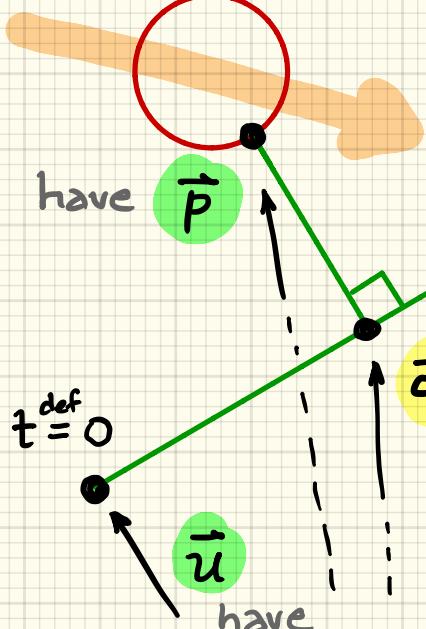
These are fundamental concepts

⇒ Apply them at CPU & GPU levels

# Basic physics calcs: wall vs. puck

$$t \stackrel{\text{def}}{=} 1$$

## COLLINEAR POINT AND PARAMETER



$$t = \frac{(\vec{P} - \vec{u}) \cdot (\vec{n} - \vec{u})}{\|\vec{n} - \vec{u}\|^2}$$

$$\vec{q} = \vec{u} + t(\vec{n} - \vec{u})$$

---

1. Find contact normal vector

$$\frac{\vec{q} - \vec{c}}{|\vec{q} - \vec{c}|} \stackrel{\text{def}}{=} \hat{n} \quad \vec{c} \stackrel{\text{def}}{=} \text{center of puck}$$

---

2. Find tangent vector

$$\begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \times \begin{bmatrix} \hat{n}_x \\ \hat{n}_y \\ 0 \end{bmatrix} = \begin{bmatrix} \hat{n}_y \\ -\hat{n}_x \\ 0 \end{bmatrix} \stackrel{\text{def}}{=} \hat{t}$$

---

3. Resolve velocity

$$\vec{v} = (\vec{v} \cdot \hat{n}) \hat{n} + (\vec{v} \cdot \hat{t}) \hat{t}$$

---

4. Reverse normal component :

$$\vec{v}' = -(\vec{v} \cdot \hat{n}) \hat{n} + (\vec{v} \cdot \hat{t}) \hat{t}$$

---

Wall is an infinite source of momentum  
=> because its mass is infinite

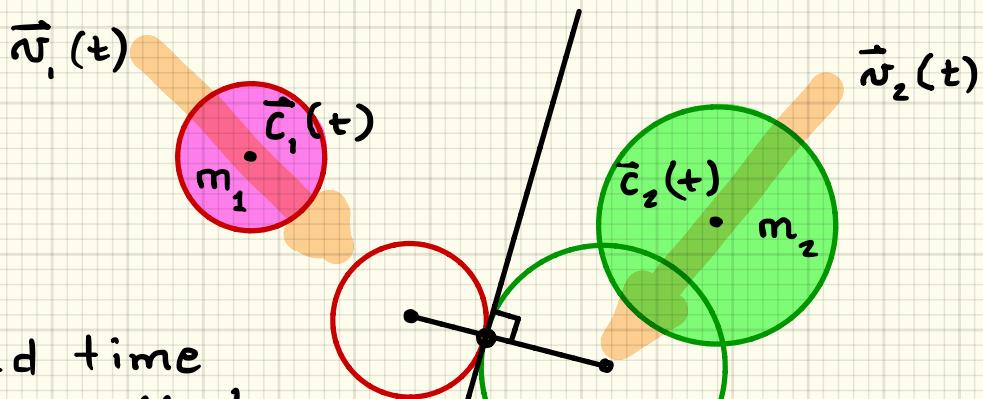
Puck energy is conserved

$$\frac{1}{2} m \vec{v} \cdot \vec{v} = \frac{1}{2} m \vec{v}' \cdot \vec{v}'$$

(check this at run time)

=> collision is perfectly elastic

## Basic Physics calcs : puck vs. puck



find time  
 $\tau$  such that

$$\vec{d}(\tau) \cdot \vec{d}(\tau) = (r_1 + r_2)^2$$

$$\begin{aligned}\vec{d}(\tau) &\stackrel{\text{def}}{=} \vec{c}_2(\tau) - \vec{c}_1(\tau) \\ &= \vec{c}_2(\text{now}) - \vec{n}_2(\text{now})\tau \\ &\quad - \vec{c}_1(\text{now}) + \vec{n}_1(\text{now})\tau\end{aligned}$$

$\Leftrightarrow$  find the zero of  $t$  the zero of  $\vec{v}$

you do the arithmetic

$$\underbrace{\|\vec{v}\|^2 t^2 - 2 \vec{d}(\text{now}) \cdot \vec{v} t + \|\vec{d}(\text{now})\|^2 - d_1^2}_{\vec{v} \stackrel{\text{def}}{=} \vec{v}_2(\text{now}) - \vec{v}_1(\text{now})} = 0$$

$$d_1 \stackrel{\text{def}}{=} r_1 + r_2$$

$$a \stackrel{\text{def}}{=} \|\vec{v}\|^2$$

$$b \stackrel{\text{def}}{=} -2 \vec{d}(\text{now}) \cdot \vec{v}$$

$$c \stackrel{\text{def}}{=} \|\vec{d}(\text{now})\|^2 - d_1^2$$

$$t = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$

collision at smallest positive real root  
if  $\exists$

Conservation of  
momentum & energy implies

(you do the  
arithmetic)

$$\vec{v}'_{1n} = \left[ (m_1 - m_2) \vec{v}_{1n} + 2m_2 \vec{v}_{2n} \right] / M$$

$$\vec{v}'_{2n} = \left[ (m_2 - m_1) \vec{v}_{2n} + 2m_1 \vec{v}_{1n} \right] / M$$

where subscript "n" means "normal component." Tangential components not changed.  $M \stackrel{\text{def}}{=} m_1 + m_2$

Check conservation every ~~step~~<sup>event</sup>: scalar equation

$$\frac{1}{2}m_1\vec{v}_1'^2 + \frac{1}{2}m_2\vec{v}_2'^2 = \frac{1}{2}m_1\vec{v}_1^2 + \frac{1}{2}m_2\vec{v}_2^2$$

$$m_1\vec{v}_1' + m_2\vec{v}_2' = m_1\vec{v}_1 + m_2\vec{v}_2$$
 vector equation

Extend the system with

- 1. rotation (spin)
- 2. puck-table friction
- 3. puck-puck friction
- 4. inelasticity
- 5. jumps off the table

} Can still be solved exactly  
i.e., pool-ball  
discrete-event sim

## INVITATION 1 (there are more)

Improve my demo app:

- desymmetrize the pucks: each schedules each collision for the other. That's double work for puck-puck collisions
- add more pucks
- add more instrumentation, strengthen the verification, find bugs
- go to 3 dims & friction: pucks → pool

## INVITATION 2:

make this run on multiple processors  
(hard, but not too hard. Try MPI. You  
will have to implement lots more of TW)

Measure additional speedup due to  
parallelism

(my one-processor demo only shows  
speedup from sequential prediction, though  
it also shows rollback & cancellation)

## INVITATION 3:

parallelize the table

let sectors coordinate ownership & collisions  
(very hard: a puck/ball can be in as many as four sectors at once).

Add many more pucks

## INVITATION 4:

make this a full O.S.

(open problem; hasn't been done since '84)

all other known impls are on TOP of  
other host O.S.s (ROSS, GoTW)

do it on a raw VM

Why is this good? ...

	Ordinary OS	TWOS
scheduling	round-robin, ...	lowest lvt first
synchronization	lock, semaphore mutex, pipe, ...	rollback
flowcontrol	throttling, windowing	cancelback
commitment	transactions	GVT
memory mgt	garbage collection, page fault	fossil collection (also GVT) rollback (time fault)
load mgt	<u>balancing</u>	process migration by LVT

Every OS function has a VT analog.

If TW is NOT an OS, all overhead is duplicated.

## INVITATION 5:

distributed fault tolerance

MIX TW with PAXOS / RAFT ?

Never been studied or tried AFAIK

Virtual Time will be a reliable guide

Multiple copies of a process can execute  
in parallel with arbitrary skew in VT

Rollback & Cancellation ensure correctness

Time Warp is inherently Byzantine!

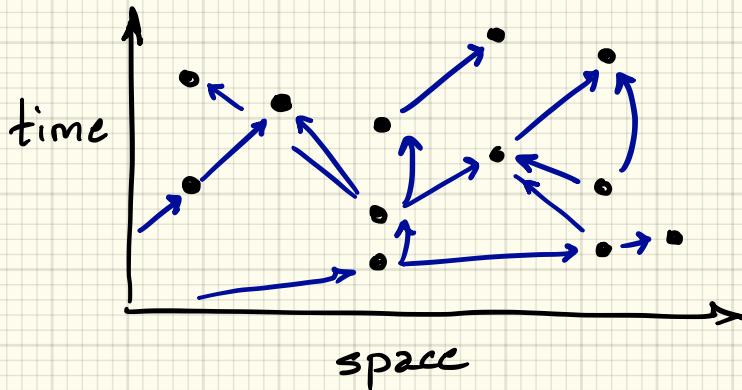
## INVITATION 6 :

full space-time programming model

Events as points in spacetime

Causality emergent rather than forced

Function calls have space AND time loci



INVITATION 7:

SIMULATE AMAZON! / on AWS  
of course!

run sims in seconds instead of weeks!

six - pager ?

AMZN on AMZN

## WRAP - UP

TW is best known method for parallel,  
distributed discrete-event simulation

Speedups of  $> 1M$  have been demonstrated

It's been implemented many times, but  
not enough times as a full OS

There are many juicy research problems

Its symmetries often make the  
impossible relatively easy